

Overlap pays with offload weight — but is capped by GPU throughput

